

# Assignment D

```
library(nycflights13)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(ggplot2)
```

```
#1 Had an arrival delay of two or more hours
flights %>%
  filter(arr_delay >= 120)
```

# A tibble: 10,200 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	811	630	101	1047	830
2	2013	1	1	848	1835	853	1001	1950
3	2013	1	1	957	733	144	1056	853
4	2013	1	1	1114	900	134	1447	1222
5	2013	1	1	1505	1310	115	1638	1431
6	2013	1	1	1525	1340	105	1831	1626
7	2013	1	1	1549	1445	64	1912	1656
8	2013	1	1	1558	1359	119	1718	1515
9	2013	1	1	1732	1630	62	2028	1825
10	2013	1	1	1803	1620	103	2008	1750

# i 10,190 more rows

# i 11 more variables: arr\_delay <dbl>, carrier <chr>, flight <int>,  
# tailnum <chr>, origin <chr>, dest <chr>, air\_time <dbl>, distance <dbl>,  
# hour <dbl>, minute <dbl>, time\_hour <dtm>

```
#2 Flew to Houston (IAH or HOU)
flights %>%
  filter(dest %in% c("IAH", "HOU"))
```

# A tibble: 9,313 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>

1	2013	1	1	517	515	2	830	819
2	2013	1	1	533	529	4	850	830
3	2013	1	1	623	627	-4	933	932
4	2013	1	1	728	732	-4	1041	1038
5	2013	1	1	739	739	0	1104	1038
6	2013	1	1	908	908	0	1228	1219
7	2013	1	1	1028	1026	2	1350	1339
8	2013	1	1	1044	1045	-1	1352	1351
9	2013	1	1	1114	900	134	1447	1222
10	2013	1	1	1205	1200	5	1503	1505

# i 9,303 more rows

# i 11 more variables: arr\_delay <dbl>, carrier <chr>, flight <int>,  
# tailnum <chr>, origin <chr>, dest <chr>, air\_time <dbl>, distance <dbl>,  
# hour <dbl>, minute <dbl>, time\_hour <dtm>

```
#3 Were operated by United, American, or Delta
flights %>%
  filter(carrier %in% c("UA", "AA", "DL"))
```

# A tibble: 139,504 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	517	515	2	830	819
2	2013	1	1	533	529	4	850	830
3	2013	1	1	542	540	2	923	850
4	2013	1	1	554	600	-6	812	837
5	2013	1	1	554	558	-4	740	728
6	2013	1	1	558	600	-2	753	745
7	2013	1	1	558	600	-2	924	917
8	2013	1	1	558	600	-2	923	937
9	2013	1	1	559	600	-1	941	910
10	2013	1	1	559	600	-1	854	902

# i 139,494 more rows

# i 11 more variables: arr\_delay <dbl>, carrier <chr>, flight <int>,  
# tailnum <chr>, origin <chr>, dest <chr>, air\_time <dbl>, distance <dbl>,  
# hour <dbl>, minute <dbl>, time\_hour <dtm>

```
#4 Departed in summer (July, August, and September)
flights %>%
  filter(month %in% c(7, 8, 9))
```

# A tibble: 86,326 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	7	1	1	2029	212	236	2359
2	2013	7	1	2	2359	3	344	344
3	2013	7	1	29	2245	104	151	1
4	2013	7	1	43	2130	193	322	14
5	2013	7	1	44	2150	174	300	100

6	2013	7	1	46	2051	235	304	2358
7	2013	7	1	48	2001	287	308	2305
8	2013	7	1	58	2155	183	335	43
9	2013	7	1	100	2146	194	327	30
10	2013	7	1	100	2245	135	337	135

# i 86,316 more rows

# i 11 more variables: arr\_delay <dbl>, carrier <chr>, flight <int>,  
# tailnum <chr>, origin <chr>, dest <chr>, air\_time <dbl>, distance <dbl>,  
# hour <dbl>, minute <dbl>, time\_hour <dtm>

```
#5 Arrived more than two hours late but didn't leave late
flights %>%
  filter(arr_delay > 120, dep_delay <= 0)
```

# A tibble: 29 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	27	1419	1420	-1	1754	1550
2	2013	10	7	1350	1350	0	1736	1526
3	2013	10	7	1357	1359	-2	1858	1654
4	2013	10	16	657	700	-3	1258	1056
5	2013	11	1	658	700	-2	1329	1015
6	2013	3	18	1844	1847	-3	39	2219
7	2013	4	17	1635	1640	-5	2049	1845
8	2013	4	18	558	600	-2	1149	850
9	2013	4	18	655	700	-5	1213	950
10	2013	5	22	1827	1830	-3	2217	2010

# i 19 more rows

# i 11 more variables: arr\_delay <dbl>, carrier <chr>, flight <int>,  
# tailnum <chr>, origin <chr>, dest <chr>, air\_time <dbl>, distance <dbl>,  
# hour <dbl>, minute <dbl>, time\_hour <dtm>

```
#6 Were delayed by at least an hour, but made up over 30 minutes in flight
flights %>%
  filter(dep_delay >= 60, (dep_delay - arr_delay) > 30)
```

# A tibble: 1,844 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	2205	1720	285	46	2040
2	2013	1	1	2326	2130	116	131	18
3	2013	1	3	1503	1221	162	1803	1555
4	2013	1	3	1839	1700	99	2056	1950
5	2013	1	3	1850	1745	65	2148	2120
6	2013	1	3	1941	1759	102	2246	2139
7	2013	1	3	1950	1845	65	2228	2227
8	2013	1	3	2015	1915	60	2135	2111
9	2013	1	3	2257	2000	177	45	2224
10	2013	1	4	1917	1700	137	2135	1950

```
# i 1,834 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

## problem 4

---

```
library(nycflights13)
library(dplyr)
library(ggplot2)

# Count the number of flights per day
daily_flights <- flights %>%
  group_by(year, month, day) %>%
  summarise(n = n(), .groups = "drop")

# Check if there are any days with zero flights
any_zero_days <- any(daily_flights$n == 0)
any_zero_days # FALSE means there was at least one flight every day
```


```
[1] FALSE
```

```
# List the days with zero flights (if any)
daily_flights %>%
  filter(n == 0)
```

```
# A tibble: 0 × 4
# i 4 variables: year <int>, month <int>, day <int>, n <int>
```

```
# Visualization: number of flights per day in 2013
ggplot(daily_flights, aes(x = as.Date(sprintf("%d-%02d-%02d", year, month, day))),
  geom_line(color = "steelblue") +
  labs(title = "Number of NYC Flights per Day in 2013",
    x = "Date",
    y = "Number of Flights") +
  theme_minimal())
```

## Number of NYC Flights per Day in 2013



## problem 5

```
# Flight(s) with the farthest distance
farthest <- flights %>%
  filter(distance == max(distance, na.rm = TRUE))

# Flight(s) with the shortest distance
shortest <- flights %>%
  filter(distance == min(distance, na.rm = TRUE))

farthest
```

# A tibble: 342 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	857	900	-3	1516	1530
2	2013	1	2	909	900	9	1525	1530
3	2013	1	3	914	900	14	1504	1530
4	2013	1	4	900	900	0	1516	1530
5	2013	1	5	858	900	-2	1519	1530
6	2013	1	6	1019	900	79	1558	1530
7	2013	1	7	1042	900	102	1620	1530

```

8 2013      1      8      901      900      1      1504      1530
9 2013      1      9      641      900     1301      1242      1530
10 2013     1     10      859      900     -1      1449      1530
# i 332 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>

```

shortest

```

# A tibble: 1 × 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>      <dbl>      <int>         <int>
1  2013     7    27      NA             106        NA        NA             245
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>

```

## Exercise 3.3.5

### problem 1

```

hm_to_min <- function(x) {
  hour <- x %/% 100
  minute <- x %% 100
  hour * 60 + minute
}

flights_check <- flights %>%
  filter(!is.na(dep_time), !is.na(sched_dep_time)) %>%
  mutate(
    dep_time_min = hm_to_min(dep_time),
    sched_dep_time_min = hm_to_min(sched_dep_time),
    recomputed_delay = dep_time_min - sched_dep_time_min,
    # fix overnight issue: if recomputed < -1000, add 1440 minutes
    recomputed_delay = if_else(recomputed_delay < -1000,
                              recomputed_delay + 1440,
                              recomputed_delay),
    diff = recomputed_delay - dep_delay
  )

summary(flights_check$diff)

```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-1440.0000	0.0000	0.0000	-0.1972	0.0000	0.0000

## problem 4

```
variables <- c("year", "month", "day", "dep_delay", "arr_delay")
```

```
# Example 1: using any_of() – safe, skips missing vars
```

```
flights %>%  
  select(any_of(variables)) %>%  
  head()
```

```
# A tibble: 6 × 5
```

	year	month	day	dep_delay	arr_delay
	<int>	<int>	<int>	<dbl>	<dbl>
1	2013	1	1	2	11
2	2013	1	1	4	20
3	2013	1	1	2	33
4	2013	1	1	-1	-18
5	2013	1	1	-6	-25
6	2013	1	1	-4	12

```
# Example 2: what happens if a variable name is wrong
```

```
bad_vars <- c("year", "month", "day", "arr_dlay") # typo: arr_delay -> arr_dlay
```

```
# any_of(): will quietly skip the missing column
```

```
flights %>%  
  select(any_of(bad_vars)) %>%  
  head()
```

```
# A tibble: 6 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1

```
# all_of(): will throw an error if any column is missing
```

## Exercise3.3.7

### problem 1

```
library(nycflights13)  
library(dplyr)
```

```
library(ggplot2)

# 1. Average arrival delay by carrier
avg_delay <- flights %>%
  group_by(carrier) %>%
  summarise(
    mean_arr_delay = mean(arr_delay, na.rm = TRUE),
    n = n(),
    .groups = "drop"
  ) %>%
  arrange(desc(mean_arr_delay))

print(avg_delay)
```

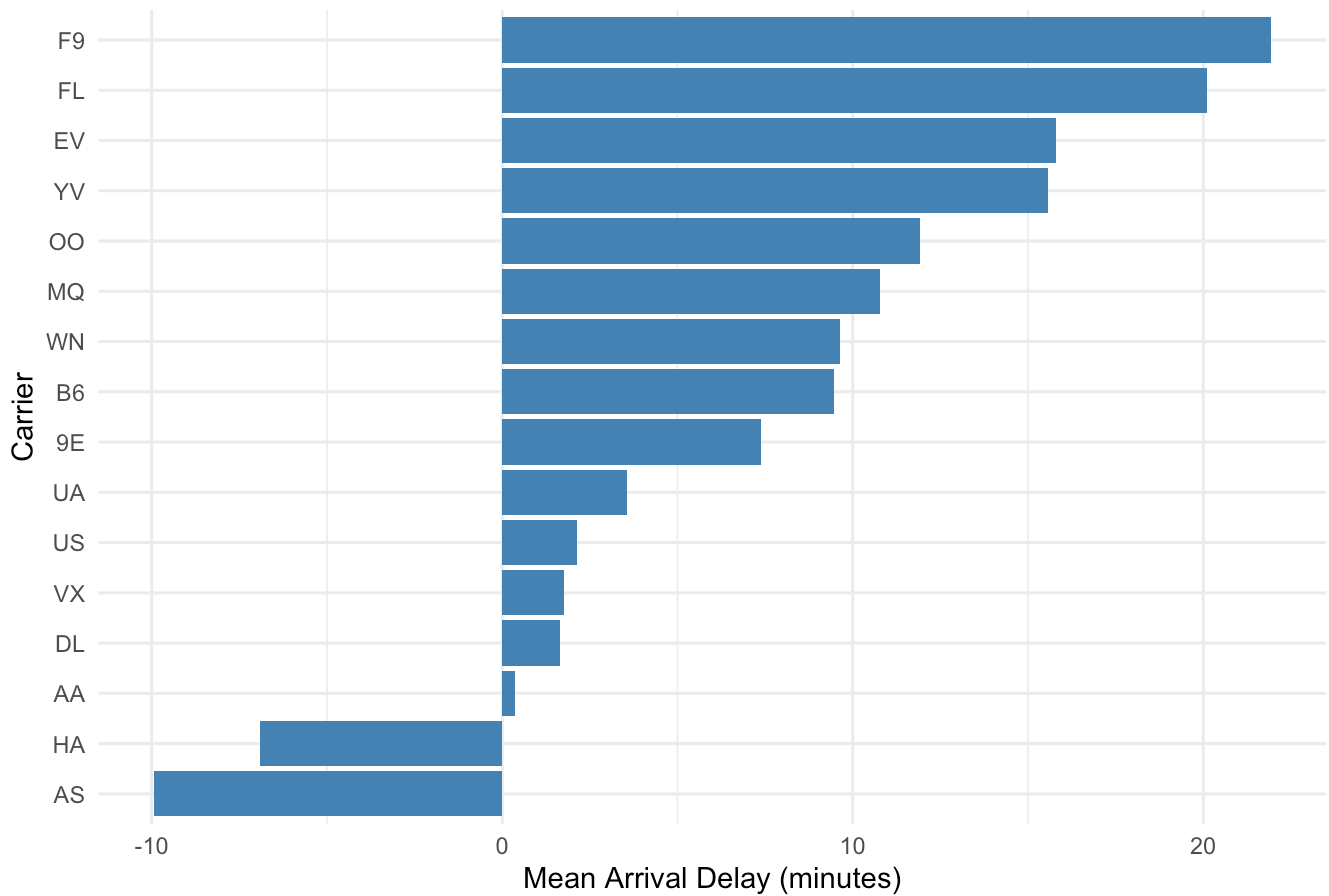
# A tibble: 16 × 3

	carrier	mean_arr_delay	n
	<chr>	<dbl>	<int>
1	F9	21.9	685
2	FL	20.1	3260
3	EV	15.8	54173
4	YV	15.6	601
5	OO	11.9	32
6	MQ	10.8	26397
7	WN	9.65	12275
8	B6	9.46	54635
9	9E	7.38	18460
10	UA	3.56	58665
11	US	2.13	20536
12	VX	1.76	5162
13	DL	1.64	48110
14	AA	0.364	32729
15	HA	-6.92	342
16	AS	-9.93	714

```
# Plot: average delay by carrier
ggplot(avg_delay, aes(x = reorder(carrier, mean_arr_delay), y = mean_arr_delay))
  geom_col(fill = "steelblue") +
  coord_flip() +
  labs(title = "Average Arrival Delay by Carrier (NYC 2013)",
       x = "Carrier",
       y = "Mean Arrival Delay (minutes)") +
  theme_minimal()
```



Average Arrival Delay by Carrier (NYC 2013)



```
# 2. Carrier x destination delays
carrier_dest_delay <- flights %>%
  group_by(carrier, dest) %>%
  summarise(
    mean_delay = mean(arr_delay, na.rm = TRUE),
    n = n(),
    .groups = "drop"
  )

# Weighted average (accounting for # of flights per destination)
weighted_delays <- carrier_dest_delay %>%
  group_by(carrier) %>%
  summarise(weighted_mean_delay = weighted.mean(mean_delay, n), .groups = "drop")
  arrange(desc(weighted_mean_delay))

print(weighted_delays)
```

```
# A tibble: 16 × 2
  carrier weighted_mean_delay
  <chr>          <dbl>
1 F9             21.9
2 FL             20.1
3 EV             15.8
```

4	YV	15.6
5	00	11.3
6	MQ	10.8
7	WN	9.65
8	B6	9.46
9	UA	3.57
10	VX	1.77
11	DL	1.65
12	AA	0.359
13	HA	-6.92
14	AS	-9.93
15	9E	NaN
16	US	NaN

```
# 3. Compare carriers at the same destination
# This shows whether a "bad carrier" is still worse at the same airport
dest_carrier_compare <- flights %>%
  group_by(dest, carrier) %>%
  summarise(mean_delay = mean(arr_delay, na.rm = TRUE),
            n = n(),
            .groups = "drop") %>%
  arrange(dest, desc(mean_delay))

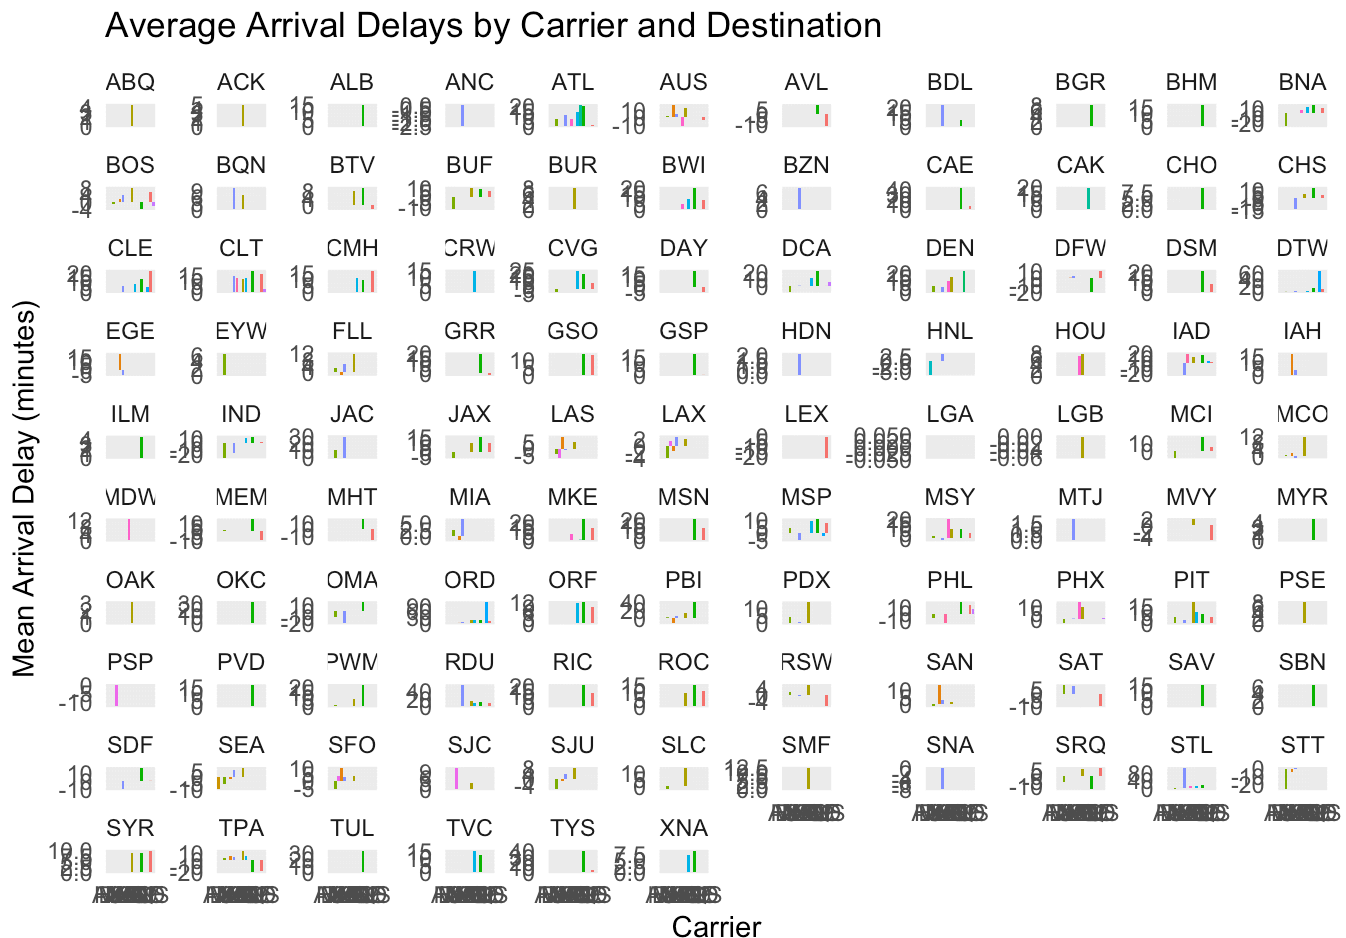
head(dest_carrier_compare, 20) # show top few
```

# A tibble: 20 × 4

	dest <chr>	carrier <chr>	mean_delay <dbl>	n <int>
1	ABQ	B6	4.38	254
2	ACK	B6	4.85	265
3	ALB	EV	14.4	439
4	ANC	UA	-2.5	8
5	ATL	FL	20.7	2337
6	ATL	EV	19.6	1764
7	ATL	MQ	14.0	2322
8	ATL	UA	10.5	103
9	ATL	DL	7.42	10571
10	ATL	WN	6.90	59
11	ATL	9E	0.857	59
12	AUS	AA	16.2	365
13	AUS	B6	11.7	747
14	AUS	UA	4.28	670
15	AUS	DL	1.41	357
16	AUS	9E	-3.5	2
17	AUS	WN	-11.2	298
18	AVL	EV	8.80	265
19	AVL	9E	-12.1	10
20	BDL	UA	22.6	8

```
# 4. Visualization: mean delay by carrier, faceted by destination
ggplot(dest_carrier_compare, aes(x = reorder(carrier, mean_delay), y = mean_delay,
  geom_col() +
  facet_wrap(~ dest, scales = "free_y") +
  theme_minimal() +
  theme(legend.position = "none") +
  labs(title = "Average Arrival Delays by Carrier and Destination",
    x = "Carrier",
    y = "Mean Arrival Delay (minutes)"))
```

Warning: Removed 2 rows containing missing values or values outside the scale range (`geom\_col()`).



## problem 2

```
# Find the most delayed departure for each destination
worst_dep_by_dest <- flights %>%
  filter(!is.na(dep_delay)) %>%
  group_by(dest) %>%
  slice_max(order_by = dep_delay, n = 1, with_ties = FALSE) %>%
  ungroup() %>%
```

```

select(year, month, day, carrier, flight, origin, dest,
       sched_dep_time, dep_time, dep_delay, arr_delay)

print(worst_dep_by_dest)

```

# A tibble: 104 × 11

	year	month	day	carrier	flight	origin	dest	sched_dep_time	dep_time
	<int>	<int>	<int>	<chr>	<int>	<chr>	<chr>	<int>	<int>
1	2013	12	14	B6	65	JFK	ABQ	2001	2223
2	2013	7	23	B6	1491	JFK	ACK	800	1139
3	2013	1	25	EV	4309	EWR	ALB	2000	123
4	2013	8	17	UA	887	EWR	ANC	1625	1740
5	2013	7	22	DL	2047	LGA	ATL	759	2257
6	2013	7	10	UA	503	EWR	AUS	1505	2056
7	2013	6	14	EV	4519	EWR	AVL	816	1158
8	2013	2	21	EV	4103	EWR	BDL	1316	1728
9	2013	12	1	EV	5309	LGA	BGR	1056	1504
10	2013	4	10	EV	5038	LGA	BHM	1900	25

# i 94 more rows

# i 2 more variables: dep\_delay <dbl>, arr\_delay <dbl>

## problem4

```

df <- tibble(x = c(5, 2, 8, 1, 9))

# 1. Positive n with slice_min(): pick the 2 smallest values
df %>% slice_min(x, n = 2)

```

# A tibble: 2 × 1

	x
	<dbl>
1	1
2	2

# Expected: rows with x = 1, 2

```

# 2. Negative n with slice_min(): drop the 2 smallest values
df %>% slice_min(x, n = -2)

```

# A tibble: 3 × 1

	x
	<dbl>
1	1
2	2
3	5

# Expected: rows with x = 5, 8, 9

```
# 3. Positive n with slice_max(): pick the 2 largest values
df %>% slice_max(x, n = 2)
```

```
# A tibble: 2 × 1
```

```
      x
<dbl>
1     9
2     8
```

```
# Expected: rows with x = 9, 8
```

```
# 4. Negative n with slice_max(): drop the 2 largest values
df %>% slice_max(x, n = -2)
```

```
# A tibble: 3 × 1
```

```
      x
<dbl>
1     9
2     8
3     5
```

```
# Expected: rows with x = 5, 2, 1
```

```
# Positive n: find the 5 most delayed departures
```

```
flights %>%
  filter(!is.na(dep_delay)) %>%
  slice_max(dep_delay, n = 5) %>%
  select(year, month, day, carrier, flight, origin, dest,
         sched_dep_time, dep_time, dep_delay)
```

```
# A tibble: 5 × 10
```

	year	month	day	carrier	flight	origin	dest	sched_dep_time	dep_time
	<int>	<int>	<int>	<chr>	<int>	<chr>	<chr>	<int>	<int>
1	2013	1	9	HA	51	JFK	HNL	900	641
2	2013	6	15	MQ	3535	JFK	CMH	1935	1432
3	2013	1	10	MQ	3695	EWR	ORD	1635	1121
4	2013	9	20	AA	177	JFK	SFO	1845	1139
5	2013	7	22	MQ	3075	JFK	CVG	1600	845

```
# i 1 more variable: dep_delay <dbl>
```

```
# Negative n: drop the 5 most delayed departures, keep the rest
```

```
flights %>%
  filter(!is.na(dep_delay)) %>%
  slice_max(dep_delay, n = -5) %>%
  summarise(total_remaining = n())
```

```
# A tibble: 1 × 1
```

```
total_remaining
```

```
1      <int>  
      328516
```

## Problem 6

---

**a**

```
df <- tibble(  
  x = 1:5,  
  y = c("a", "b", "a", "a", "b"),  
  z = c("K", "K", "L", "L", "K")  
)  
df
```

```
# A tibble: 5 × 3  
      x y      z  
  <int> <chr> <chr>  
1     1 a      K  
2     2 b      K  
3     3 a      L  
4     4 a      L  
5     5 b      K
```

```
df |>  
  group_by(y)
```

```
# A tibble: 5 × 3  
# Groups:   y [2]  
      x y      z  
  <int> <chr> <chr>  
1     1 a      K  
2     2 b      K  
3     3 a      L  
4     4 a      L  
5     5 b      K
```

so the data will be the same, but grouping information is added.

**b**

```
df <- tibble(  
  x = 1:5,  
  y = c("a", "b", "a", "a", "b"),  
  z = c("K", "K", "L", "L", "K")  
)  
  
df
```

```
# A tibble: 5 × 3
  x y     z
<int> <chr> <chr>
1     1 a     K
2     2 b     K
3     3 a     L
4     4 a     L
5     5 b     K
```

```
df |> arrange(y)
```

```
# A tibble: 5 × 3
  x y     z
<int> <chr> <chr>
1     1 a     K
2     3 a     L
3     4 a     L
4     2 b     K
5     5 b     K
```

Arrange() reorders the rows of the data frame according to the values of one or more columns.

## C

We expect to see a 2x2 table.

```
df |>
  group_by(y) |>
  summarize(mean_x = mean(x))
```

```
# A tibble: 2 × 2
  y     mean_x
<chr> <dbl>
1 a       2.67
2 b       3.5
```

group\_by(y) tells R to treat rows with the same value of y as belonging to the same group, and summarize(mean\_x = mean(x)) then calculates the mean of x within each group. The result is a collapsed summary table that contains one row per group along with the group labels and their corresponding summary statistics.

## D

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x))
```

``summarise()`` has grouped output by 'y'. You can override using the `` .groups`` argument.

```
# A tibble: 3 × 3
# Groups:   y [2]
  y     z   mean_x
  <chr> <chr> <dbl>
1 a     K     1
2 a     L     3.5
3 b     K     3.5
```

`group_by(y, z)` divides the data into groups defined by each unique combination of y and z, and `summarize(mean_x = mean(x))` then calculates the mean of x within each group. The result is a summary tibble where each row corresponds to a unique (y, z) pair along with its calculated group statistic.

## E

```
df |>
  group_by(y, z) |>
  summarize(mean_x = mean(x), .groups = "drop")
```

```
# A tibble: 3 × 3
  y     z   mean_x
  <chr> <chr> <dbl>
1 a     K     1
2 a     L     3.5
3 b     K     3.5
```

`group_by(y, z)` creates subgroups based on each unique combination of y and z, and `summarize(mean_x = mean(x), .groups = "drop")` computes the mean of x within each subgroup while removing all grouping information from the result. The final output is a simple tibble with one row per (y, z) pair and no residual grouping.

## F

```
df |>
  group_by(y, z) |>
  mutate(mean_x = mean(x))
```

```
# A tibble: 5 × 4
# Groups:   y, z [3]
  x y     z   mean_x
  <int> <chr> <chr> <dbl>
1 1 a     K     1
2 2 b     K     3.5
3 3 a     L     3.5
4 4 a     L     3.5
5 5 b     K     3.5
```



The `summarize()` pipeline collapses each group into a single row, producing a smaller tibble with one row per (y, z) combination, while the `mutate()` pipeline keeps the original number of rows and simply adds a new column containing the group's mean repeated across all rows in that group. In other words, `summarize()` reduces the data, whereas `mutate()` augments it by attaching group statistics without changing the row count.